

FORECASTING POWER CONSUMPTION

A PROJECT REPORT

submitted by

Pranti Debnath

debnathpranti@gmail.com

to

Shehan Irteza

prantos@uab.edu

&

Rahad Arman Nabid

rahad.arman.nabid@temple.edu



Department of Electronics and Telecommunication Engineering

Chittagong University of Engineering and Technology

February 2025

CONTENTS

LIST OF TABLES	3
Chapter 1. UNDERSTANDING THE DATASET	1
1.1 Key Features and Their Relevance to Forecasting	1
1.2 Challenges in Forecasting Time-Series Data	2
1.3 Impact of Weather on Power Consumption	3
Chapter 2. DATA PREPROCESSING	5
2.1 Handling Missing or Inconsistent Data	5
2.2 Avoiding Feature Engineering	5
2.3 Splitting the Dataset	5
2.4 Tokenizing Time-Series Data	6
Chapter 3. MODEL IMPLEMENTATION	7
3.1 Implementing Forecasting Models	7
3.2 Comparing Model Architectures	7
Chapter 4. FINE-TUNING EACH MODEL	8
4.1 Experimenting with Hyperparameters	8
4.2 Regularization Techniques	9
4.3 Optimizers and Learning Rate Scheduling	9
4.4 Early Stopping Implementation	10
Chapter 5. IMPLEMENTATION	11
5.1 Evaluation Metrics Used	11
5.2 Implementing Evaluation Metrics	11
5.3 Visualization Techniques	12

5.4 Best Performing Model	13
Chapter 6. CONCLUSION	14

LIST OF TABLES

3.1 Comparing Model Architectures	7
---	---

Chapter 1

UNDERSTANDING THE DATASET

1.1 KEY FEATURES AND THEIR RELEVANCE TO FORECASTING

The dataset comprises a variety of features that directly or indirectly influence the ability to forecast power consumption. Each feature plays a significant role in shaping the predictive models:

- **Datetime:** This serves as the backbone of the dataset, acting as the temporal index. The Datetime feature captures crucial temporal dynamics, including hourly, daily, and seasonal variations that are integral to understanding and predicting consumption trends.
- **Temperature:** One of the most influential features, temperature directly correlates with energy consumption. High temperatures typically lead to increased usage of cooling systems, whereas low temperatures increase heating requirements. The variability of temperature across seasons makes it a critical feature for forecasting.
- **Humidity:** This feature indirectly impacts energy usage by influencing HVAC (Heating, Ventilation, and Air Conditioning) systems. Higher humidity levels often lead to increased demand for dehumidification or air conditioning.
- **Wind Speed:** While the direct impact of wind speed on energy consumption may be less pronounced, it provides environmental context. Windier

conditions might impact outdoor heating or cooling demands and indirectly influence energy usage patterns.

- **General Diffuse Flows and Diffuse Flows:** These are indicators of solar radiation, which directly impact the availability of renewable energy resources. In regions heavily reliant on solar power, these variables can provide critical insights into energy availability and demand balance.
- **PowerConsumption Zone1, Zone2, Zone3:** These target variables represent the energy consumption levels in three distinct zones. Forecasting these variables accurately is the primary objective of the analysis.

Using SHAP (SHapley Additive exPlanations) analysis, the relative importance of these features is determined. SHAP values highlight how much each feature contributes to the model's prediction for each individual instance. For instance, temperature and humidity frequently emerge as top contributors, especially during extreme weather conditions.

1.2 CHALLENGES IN FORECASTING TIME-SERIES DATA

Forecasting time-series data presents unique challenges that require careful consideration and robust methodologies:

- **Seasonality and Trends:** Energy consumption data often exhibit strong seasonality and underlying trends. Capturing daily, weekly, and annual cycles accurately is essential for robust forecasting.
- **Data Noise:** Noise in the data, introduced by unexpected events such as public holidays, maintenance activities, or extreme weather conditions, can distort patterns and reduce model accuracy.

- **Long-Range Dependencies:** In time-series forecasting, long-term dependencies where past events continue to influence future observations must be accounted for. This is particularly challenging for simpler models and necessitates advanced architectures like transformers.
- **Multicollinearity:** Weather-related features often exhibit high correlations among themselves. For example, temperature and humidity might co-vary. This multicollinearity can confuse simpler models and necessitates advanced techniques to disentangle relationships.

Exploratory data analysis of the dataset reveals significant differences in weekday and weekend consumption patterns, as well as daily seasonality. This underscores the need for models that can effectively capture such temporal dynamics.

1.3 IMPACT OF WEATHER ON POWER CONSUMPTION

Weather-related variables such as temperature, humidity, and wind speed demonstrate varying levels of correlation with power consumption across the three zones. Their impact can be summarized as follows:

- **Zone 1:** Temperature exhibits a strong positive correlation with power consumption ($\rho = 0.85$). This indicates that heating and cooling demands in this zone are highly sensitive to temperature fluctuations.
- **Zone 2:** Humidity demonstrates a moderate negative correlation ($\rho = -0.43$), suggesting that energy demands are influenced by both humid and dry conditions. The interplay between humidity and HVAC efficiency is particularly notable in this zone.
- **Zone 3:** Wind speed shows negligible correlation with power consumption. This suggests that, unlike temperature and humidity, wind speed does not play a direct role in determining energy usage in this zone.

Statistical analysis confirms the outsized influence of temperature, especially during extreme weather events. Peaks in energy demand often coincide with temperature extremes, highlighting the critical role of this feature in forecasting models. Descriptive statistics also show that zones exhibit different sensitivities to weather variables, emphasizing the need for tailored models for each zone.

Chapter 2

DATA PREPROCESSING

2.1 HANDLING MISSING OR INCONSISTENT DATA

Although the dataset contains no missing values, preprocessing involved addressing inconsistencies:

- **Outlier Handling:** Extreme values in features such as Wind Speed and Diffuse Flows were capped at the 1st and 99th percentiles to reduce noise.
- **Normalization:** Numerical columns were scaled using MinMaxScaler to ensure compatibility with models requiring normalized input.

2.2 AVOIDING FEATURE ENGINEERING

To adhere to the guidelines, no advanced feature engineering was performed. Instead, raw data was used directly. The Datetime column was parsed to extract basic time-based components such as hour, day, month, and weekday, which are essential for capturing temporal trends.

2.3 SPLITTING THE DATASET

The dataset was split into three subsets:

- **Training Set:** 70% of the data, used for model learning.
- **Validation Set:** 20%, used for hyperparameter tuning and early stopping.

- **Test Set:** 10%, reserved for final evaluation.

2.4 TOKENIZING TIME-SERIES DATA

To prepare data for transformer-based models, sequences of 96 time steps were created using a sliding window approach:

```
1 def create_sequences(data, seq_length):
2     X, y = [], []
3     for i in range(len(data) - seq_length):
4         X.append(data[i:i + seq_length, :-1])
5         y.append(data[i + seq_length, -1])
6     return np.array(X), np.array(y)
```

Chapter 3

MODEL IMPLEMENTATION

3.1 IMPLEMENTING FORECASTING MODELS

- **Vanilla Transformer:** Employs positional encoding and self-attention mechanisms to process full-length sequences, capturing short-term dependencies effectively.
- **PatchTST:** Utilizes a patch-based approach, dividing sequences into smaller patches to enhance efficiency and handle long-term dependencies more effectively.

3.2 COMPARING MODEL ARCHITECTURES

Table 3.1: Comparing Model Architectures

Feature	Vanilla Transformer	PatchTST
Sequence Handling	Processes full sequences	Processes patches
Memory Usage	High	Efficient
Performance	Better for short-term	Superior for long-term
Training Speed	Slower due to self-attention on all timesteps	Faster due to reduced sequence length

Chapter 4

FINE-TUNING EACH MODEL

4.1 EXPERIMENTING WITH HYPERPARAMETERS

Hyperparameter tuning is essential for enhancing model accuracy and computational efficiency. The following key hyperparameters were optimized:

- **Learning Rate:** Different values (0.0001, 0.001, 0.01) were tested, with 0.0001 providing the best trade-off between convergence speed and accuracy.
- **Batch Size:** A batch size of 32 was determined to be optimal for balancing memory efficiency and model performance.
- **Sequence Length:** The sequence length was set to 96, allowing the model to capture sufficient historical context while avoiding excessive computation.
- **Number of Attention Heads:** Experiments with 4 and 8 heads indicated that 8 heads provided a marginally improved ability to capture dependencies across different time steps.
- **Feedforward Dimensions:** A value of 256 was selected to maintain model complexity while ensuring effective training.
- **Dropout Rate:** A dropout rate of 0.1 was found to reduce overfitting without significantly affecting learning capacity.
- **Weight Decay:** Regularization through AdamW optimizer helped stabilize the model training process.

4.2 REGULARIZATION TECHNIQUES

Regularization was employed to prevent overfitting and enhance the model's generalization ability:

- **Dropout Regularization:** Applied with a rate of $*0.1*$ to introduce randomness and reduce over-reliance on specific features.
- **Weight Initialization:** Xavier initialization was used to ensure a stable weight distribution and prevent exploding or vanishing gradients.
- **Early Stopping:** Implemented to halt training if the validation loss did not improve for 10 consecutive epochs, preventing unnecessary overfitting.
- **Gradient Clipping:** To avoid exploding gradients, gradients were clipped at a threshold of 1.0.
- **Batch Normalization:** Introduced within the transformer layers to stabilize activations and improve convergence speed.

4.3 OPTIMIZERS AND LEARNING RATE SCHEDULING

To ensure stable training and optimal convergence, the following optimization strategies were used:

- **AdamW Optimizer:** Chosen for its efficiency in handling large datasets, enabling faster convergence than traditional stochastic gradient descent (SGD).
- **StepLR Scheduler:** Applied to reduce the learning rate by 50% every 5 epochs, ensuring smooth learning rate decay and improved generalization.
- **Warmup Scheduler:** A warmup phase for the learning rate was introduced, allowing the model to adapt to the dataset before full optimization took place.

4.4 EARLY STOPPING IMPLEMENTATION

- **Validation-Based Stopping:** Training was halted if the validation loss did not improve for 10 consecutive epochs, preventing unnecessary computational expenses and avoiding overfitting.
- **Patience Parameter:** Set to 10 epochs, ensuring that temporary fluctuations in performance did not prematurely terminate training.
- **Best Model Checkpointing:** The best-performing model was saved based on validation loss to ensure the highest-quality results were retained.

Chapter 5

IMPLEMENTATION

5.1 EVALUATION METRICS USED

Model performance was evaluated using the following key metrics:

- **Mean Absolute Error (MAE):** Measures the absolute difference between predicted and actual values, providing an overall sense of prediction accuracy.
- **Root Mean Squared Error (RMSE):** Penalizes large errors more heavily than MAE, making it suitable for identifying substantial deviations.
- **Mean Absolute Percentage Error (MAPE):** Expresses the average prediction error as a percentage of actual values, making comparisons across different zones more interpretable.
- **Coefficient of Determination (R^2):** Measures how well the model explains variance in power consumption, with values closer to 1.0 indicating better predictive accuracy.
- **Symmetric Mean Absolute Percentage Error (SMAPE):** Used to handle cases where actual values can be close to zero, ensuring stable percentage-based evaluation.

5.2 IMPLEMENTING EVALUATION METRICS

The following Python code snippet was used to compute evaluation metrics:

```

1  from sklearn.metrics import mean_absolute_error
2  from sklearn.metrics import mean_squared_error
3  from sklearn.metrics import r2_score
4  import numpy as np
5  def evaluate_model(y_true, y_pred):
6      mae = mean_absolute_error(y_true, y_pred)
7      rmse = mean_squared_error(y_true, y_pred, squared=False)
8      mape = np.mean(np.abs((y_true - y_pred) / y_true)) * 100
9      r2 = r2_score(y_true, y_pred)
10     return {
11         "MAE": mae,
12         "RMSE": rmse,
13         "MAPE": mape,
14         "R2": r2
15     }

```

5.3 VISUALIZATION TECHNIQUES

To better understand model predictions and performance, several visualization techniques were implemented:

- **Actual vs. Predicted Power Consumption:** Line plots compared real and predicted power consumption values to assess the model's tracking ability.
- **Performance Metrics Comparison:** Bar charts provided a comparative analysis of MAE, RMSE, and MAPE across models and zones.
- **SHAP Analysis:** SHAP (SHapley Additive exPlanations) was used to interpret feature importance, explaining the contribution of each variable to model predictions.

- **Residual Plots:** Residual error distributions were plotted to analyze whether errors were randomly distributed, indicating unbiased model predictions.
- **Time-Series Decomposition:** The dataset was decomposed into trend, seasonal, and residual components to highlight how well the model captured different data properties.
- **Heatmaps:** Correlation heatmaps visualized the relationships between different features and power consumption patterns across zones.

5.4 BEST PERFORMING MODEL

After evaluating both models across various performance metrics and visualization techniques, PatchTST was found to be superior in long-term forecasting and generalization. Key performance comparisons include:

- **Zone 1:** PatchTST achieved 10% lower RMSE than the Vanilla Transformer, indicating more stable predictions.
- **Zone 2:** PatchTST had 8% lower MAPE, demonstrating better adaptability to variations in energy consumption.
- **Zone 3:** PatchTST had a 12% higher R^2 score, showing greater accuracy in capturing power consumption trends.
- **Overall Computational Efficiency:** PatchTST was computationally more efficient, reducing training time while maintaining high accuracy.

Chapter 6

CONCLUSION

This study explored and compared the Vanilla Transformer and PatchTST models for power consumption forecasting. While both models were effective, PatchTST outperformed Vanilla Transformer due to its ability to efficiently capture long-range dependencies, reducing forecasting errors significantly.

Future Improvements and Considerations:

- **Hybrid Architectures:** Combining CNNs, Transformers, and LSTMs may enhance both short-term and long-term predictive performance.
- **External Data Integration:** Incorporating economic indicators, real-time weather conditions, and smart meter readings can further improve model accuracy.
- **Alternative Transformer Models:** Exploring architectures such as Informer, Reformer, or Longformer could provide more efficient long-sequence forecasting capabilities.
- **Fine-Tuned Attention Mechanisms:** Implementing adaptive attention mechanisms may further optimize performance in fluctuating demand environments.

By implementing these improvements, power consumption forecasting can be enhanced, ultimately enabling better energy distribution, load balancing, and cost efficiency in smart grid systems.